

SWINBURNE

SWINBURNE UNIVERSITY OF TECHNOLOGY

COS30019: Introduction to Artificial Intelligence

Intelligent Agents

1

Previously ...

- What is AI?
 - Four paradigms (think vs. act, human-like vs. rationally)
 - AI in movies/science fictions (food for thought)
 - AI in the real world
 - In COS30019, we will study “Systems that act rationally”
- “Systems that do the right thing”
 - The “Right thing” is the course of action that is expected to maximize goal achievement given the available information.

2

Outline

- Agents and environments.
 - The vacuum-cleaner world
- The concept of rational behavior.
- Environments.
- Agent structure.

3

Agents and environments

- Agents include human, robots, softbots, thermostats, etc.
- The agent function maps percept sequence to actions

- An agent can perceive its own actions, but not always its effects.

$$f : P^* \rightarrow A$$

4

Agents and environments

- The agent function will internally be represented by the agent program.
- The agent program runs on the physical architecture to produce f .


5

The vacuum-cleaner world – An example

- Environment:** squares A and B
- Percepts:** [location and content] e.g. [A, Dirty]
- Actions:** left, right, suck, and no-op

6


The vacuum-cleaner world – Agent function



Percept sequence	Action
[A, Clean]	Right
[A, Dirty]	Suck
[B, Clean]	Left
[B, Dirty]	Suck
[A, Clean], [A, Clean]	Right
[A, Clean], [A, Dirty]	Suck
...	...

7

The vacuum-cleaner world – An agent program



```

procedure REFLEX-VACUUM-AGENT ((location, status)) return an action
  if status == Dirty then return Suck
  else if location == A then return Right
  else if location == B then return Left
    
```

What is the right function? Can it be implemented in a small agent program?

8

The concept of rationality

- A **rational agent** is one that does the right thing.
 - Every entry in the table is filled out correctly.
- What is the right thing?
 - Approximation: the most *successful* agent.
 - *Measure of success*?
- Performance measure should be **objective**
 - E.g. the amount of dirt cleaned within a certain time.
 - E.g. how clean the floor is.
 - ...
- Performance measure according to what is wanted in the environment instead of how the agents should behave.

9

Rationality

- What is rational at a given time depends on four things:
 - Performance measure,
 - Prior environment knowledge,
 - Actions,
 - Percept sequence to date (sensors).
- DEF: A rational agent chooses whichever action that maximizes the expected value of the performance measure given the percept sequence to date and prior environment knowledge.

10

Rationality

- Rationality \neq omniscience
 - An omniscient agent knows the actual outcome of its actions.
- Rationality \neq perfection
 - Rationality maximizes *expected* performance, while perfection maximizes *actual* performance.

11

Rationality

- The proposed definition requires:
 - Information gathering/exploration
 - To maximize future rewards
 - Learn from percepts
 - Extending prior knowledge
 - Agent autonomy
 - Compensate for incorrect prior knowledge

12

Is the vacuum cleaner agent rational?

■ Depend!

■ For example, it's rational under the following assumptions:

- ☐ Performance measure: 1 point for each clean square over 'lifetime' of 1000 steps
- ☐ 'geography' known but dirt distribution, initial position of agent not known
- ☐ Clean squares stay clean, sucking cleans squares
- ☐ Left and Right don't take agent outside environment
- ☐ Available actions: Left, Right, Suck, NoOp
- ☐ Agent knows where it is and whether that location contains dirt



13

Environments

■ To design a rational agent we must specify its **task environment**

■ PEAS description of the task environment:

- ☐ Performance
- ☐ Environment
- ☐ Actuators
- ☐ Sensors



14

Environments

■ E.g. Fully automated taxi:

☐ PEAS description of the environment:

- ☐ Performance
 - ☐ Safety, destination, profits, legality, comfort
- ☐ Environment
 - ☐ Streets/freeways, other traffic, pedestrians, weather, ...
- ☐ Actuators
 - ☐ Steering, accelerating, brake, horn, speaker/display,...
- ☐ Sensors
 - ☐ Video, sonar, speedometer, engine sensors, keyboard, GPS, ...



15

Environment types

	Crossword	Backgammon	Chess w/ clock	Taxi
Observable??				
Deterministic??				
Episodic??				
Static??				
Discrete??				
Single-agent??				



16

The game of backgammon



17

Environment types

Fully vs. partially observable: an environment is full observable when the sensors can detect all aspects that are *relevant* to the choice of action.

	Crossword	Backgammon	Chess w/ clock	Taxi
Observable??				
Deterministic??				
Episodic??				
Static??				
Discrete??				
Single-agent??				



18

Environment types



Fully vs. partially observable: an environment is full observable when the sensors can detect all aspects that are *relevant* to the choice of action.

	Crossword	Backgammon	Chess w/ clock	Taxi
Observable??	FULL	FULL	FULL	PARTIAL
Deterministic??				
Episodic??				
Static??				
Discrete??				
Single-agent??				



19

Environment types



Deterministic vs. stochastic: if the next environment state is completely determined by the current state the executed action then the environment is deterministic.

	Crossword	Backgammon	Chess w/ clock	Taxi
Observable??	FULL	FULL	FULL	PARTIAL
Deterministic??				
Episodic??				
Static??				
Discrete??				
Single-agent??				



20

Environment types



Deterministic vs. stochastic: if the next environment state is completely determined by the current state the executed action then the environment is deterministic.

	Crossword	Backgammon	Chess w/ clock	Taxi
Observable??	FULL	FULL	FULL	PARTIAL
Deterministic??	YES	NO	YES	NO
Episodic??				
Static??				
Discrete??				
Single-agent??				



21

Environment types



Episodic vs. sequential: In an episodic environment the agent's experience can be divided into atomic steps where the agents perceives and then performs A single action. The choice of action depends only on the episode itself

	Crossword	Backgammon	Chess w/ clock	Taxi
Observable??	FULL	FULL	FULL	PARTIAL
Deterministic??	YES	NO	YES	NO
Episodic??				
Static??				
Discrete??				
Single-agent??				



22

Environment types



Episodic vs. sequential: In an episodic environment the agent's experience can be divided into atomic steps where the agents perceives and then performs A single action. The choice of action depends only on the episode itself

	Crossword	Backgammon	Chess w/ clock	Taxi
Observable??	FULL	FULL	FULL	PARTIAL
Deterministic??	YES	NO	YES	NO
Episodic??	NO	NO	NO	NO
Static??				
Discrete??				
Single-agent??				



23

Environment types



Static vs. dynamic: If the environment can change while the agent is choosing an action, the environment is dynamic. Semi-dynamic if the agent's performance changes even when the environment remains the same.

	Crossword	Backgammon	Chess w/ clock	Taxi
Observable??	FULL	FULL	FULL	PARTIAL
Deterministic??	YES	NO	YES	NO
Episodic??	NO	NO	NO	NO
Static??				
Discrete??				
Single-agent??				



24

Environment types

Static vs. dynamic: If the environment can change while the agent is choosing an action, the environment is dynamic. Semi-dynamic if the agent's performance changes even when the environment remains the same.

	Crossword	Backgammon	Chess w/ clock	Taxi
Observable??	FULL	FULL	FULL	PARTIAL
Deterministic??	YES	NO	YES	NO
Episodic??	NO	NO	NO	NO
Static??	YES	YES	SEMI	NO
Discrete??				
Single-agent??				



25

Environment types

Discrete vs. continuous: This distinction can be applied to the state of the environment, the way time is handled and to the percepts/actions of the agent.

	Crossword	Backgammon	Chess w/ clock	Taxi
Observable??	FULL	FULL	FULL	PARTIAL
Deterministic??	YES	NO	YES	NO
Episodic??	NO	NO	NO	NO
Static??	YES	YES	SEMI	NO
Discrete??				
Single-agent??				



26

Environment types

Discrete vs. continuous: This distinction can be applied to the state of the environment, the way time is handled and to the percepts/actions of the agent.

	Crossword	Backgammon	Chess w/ clock	Taxi
Observable??	FULL	FULL	FULL	PARTIAL
Deterministic??	YES	NO	YES	NO
Episodic??	NO	NO	NO	NO
Static??	YES	YES	SEMI	NO
Discrete??	YES	YES	YES	NO
Single-agent??				



27

Environment types

Single vs. multi-agent: Does the environment contain other agents who are also maximizing some performance measure that depends on the current agent's actions?

	Crossword	Backgammon	Chess w/ clock	Taxi
Observable??	FULL	FULL	FULL	PARTIAL
Deterministic??	YES	NO	YES	NO
Episodic??	NO	NO	NO	NO
Static??	YES	YES	SEMI	NO
Discrete??	YES	YES	YES	NO
Single-agent??				



28

Environment types

Single vs. multi-agent: Does the environment contain other agents who are also maximizing some performance measure that depends on the current agent's actions?

	Crossword	Backgammon	Chess w/ clock	Taxi
Observable??	FULL	FULL	FULL	PARTIAL
Deterministic??	YES	NO	YES	NO
Episodic??	NO	NO	NO	NO
Static??	YES	YES	SEMI	NO
Discrete??	YES	YES	YES	NO
Single-agent??	YES	NO	NO	NO



29

Environment types

- The simplest environment is
 - ☐ Fully observable, deterministic, episodic, static, discrete and single-agent.
- Most real situations are:
 - ☐ Partially observable, stochastic, sequential, dynamic, continuous and multi-agent.



30

Agent types

- How does the inside of the agent work?
 - Agent = architecture + program
- All agents have the same skeleton:
 - Input = current percepts
 - Output = action
 - Program = manipulates input to produce output
- Note difference with agent function.



31

Agent types

Function TABLE-DRIVEN_AGENT(*percept*) returns an action

static: *percepts*, a sequence initially empty
table, a table of actions, indexed by percept sequence

append *percept* to the end of *percepts*

action ← LOOKUP(*percepts*, *table*)

return *action*

This approach is doomed to failure



32

Agent types

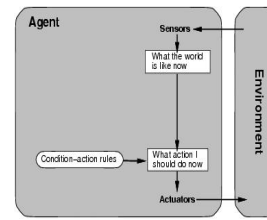
- Four basic kind of agent programs will be discussed:
 - Simple reflex agents
 - Model-based reflex agents
 - Goal-based agents
 - Utility-based agents
- All these can be turned into learning agents.
 - And that gives you four additional advanced agent types



33


Agent types; simple reflex

- Select action on the basis of *only the current percept*.
 - E.g. the vacuum-agent
- Large reduction in possible percept/action situations(next page).
- Implemented through *condition-action rules*
 - If dirty then suck



34

The vacuum-cleaner world



function REFLEX-VACUUM-AGENT (*[location, status]*) return an action
 if *status* == Dirty then return Suck
 else if *location* == A then return Right
 else if *location* == B then return Left

Reduction from 4^T to 4 entries



35

Agent types; simple reflex

function SIMPLE-REFLEX-AGENT(*percept*) returns an action

static: *rules*, a set of condition-action rules

state ← INTERPRET-INPUT(*percept*)

rule ← RULE-MATCH(*state*, *rules*)

action ← RULE-ACTION[*rule*]


return *action*

Will only work if the environment is *fully observable*
 otherwise infinite loops may occur.



36

Agent types; reflex and state



- To tackle *partially observable* environments.
 - Maintain internal state
 - Over time update state using world knowledge
 - How does the world change.
 - How do actions affect world.
- ⇒ *Model of World*



37

Agent types; reflex and state

function REFLEX-AGENT-WITH-STATE(*percept*) **returns** an action

static: *rules*, a set of condition-action rules

state, a description of the current world state

actions, the most recent actions.

state ← UPDATE-STATE(*state*, *actions*, *percept*)

rule ← RULE-MATCH(*state*, *rules*)


action ← RULE-ACTION[*rule*]

return *action*



38

Agent types; goal-based




- The agent needs a goal to know which situations are *desirable*.
 - Things become difficult when long sequences of actions are required to find the goal.
- Typically investigated in **search** and **planning** research.
- Major difference: future is taken into account
- Is more flexible since knowledge is represented explicitly and can be manipulated.



39

Agent types; utility-based




- Certain goals can be reached in different ways.
 - Some are better, have a higher utility.
- Utility function maps a (sequence of) state(s) onto a real number.
- Improves on goals:
 - Selecting between conflicting goals
 - Select appropriately between several goals based on likelihood of success.



40

Agent types; learning




- All previous agent-programs describe methods for selecting *actions*.
 - Yet it does not explain the origin of these programs.
 - Learning mechanisms can be used to perform this task.
 - Teach them instead of instructing them.
 - Advantage is the robustness of the program toward initially unknown environments.



41

Agent types; learning



- *Learning element*: introduce improvements in performance element.
 - Critic provides feedback on agents performance based on fixed performance standard.
- *Performance element*: selecting actions based on percepts.
 - Corresponds to the previous agent programs
- *Problem generator*: suggests actions that will lead to new and informative experiences.
 - Exploration vs. exploitation



42

Summary: Agents



- An **agent** perceives and acts in an environment, has an architecture, and is implemented by an agent program.
- Task environment – **PEAS** (Performance, Environment, Actuators, Sensors)
- An **ideal agent** always chooses the action which maximizes its expected performance, given its percept sequence so far.
- An **autonomous learning agent** uses its own experience rather than built-in knowledge of the environment by the designer.
- An **agent program** maps from percept to action and updates internal state.
 - **Reflex agents** respond immediately to percepts.
 - **Goal-based agents** act in order to achieve their goal(s).
 - **Utility-based agents** maximize their own utility function.
- **Representing knowledge** is important for successful agent design.
- The most challenging environments are not fully observable, nondeterministic, dynamic, and continuous

