

SWINBURNE

SWINBURNE UNIVERSITY OF TECHNOLOGY

COS30019: Introduction to Artificial Intelligence

Problem solving and Search

1

Previously on COS30019 ...

■ 4 paradigms of AI:

□ Systems that think/act like a human/rationally

■ Intelligent agents are systems that act rationally

□ chooses whichever action that maximizes the expected value of the performance measure given the percept sequence to date and prior environment knowledge

■ 4 basic types of agent & 4 (basic type + learning) agents

□ Simple reflex

□ State-based reflex

□ Goal-based agent

□ Utility-based agent

2

Outline

■ Problem-solving agents

□ A kind of goal-based agent

■ Problem types

□ Single state (fully observable)

□ Search with partial information

■ Problem formulation

□ Example problems

■ Basic search algorithms

□ Uninformed

3

Example: 8-puzzle

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

4

Example: Robot Navigation

Start state


Goal state

5

Example: Route finding

6

Example: Romania



7

Example: Romania

- On holiday in Romania; currently in Arad
 - Flight leaves tomorrow from Bucharest
- Formulate goal
 - Be in Bucharest
- Formulate problem
 - States: various cities
 - Actions: drive between cities
- Find solution
 - Sequence of cities; e.g. Arad, Sibiu, Fagaras, Bucharest, ...

8

Problem-solving agent

- Four general steps in problem solving:
 - Goal formulation
 - What are the successful world states
 - Problem formulation
 - What actions and states to consider given the goal
 - Search
 - Determine the possible sequence of actions that lead to the states of known values and then choosing the best sequence.
 - Execute
 - Given the solution, perform the actions.

9


Problem types

- Deterministic, fully observable \Rightarrow *single state problem*
 - Agent knows exactly which state it will be in; solution is a sequence.
- Partial knowledge of states and actions:
 - Non-observable \Rightarrow *sensorless or conformant problem*
 - Agent may have no idea where it is; solution (if any) is a sequence.
 - Nondeterministic and/or partially observable \Rightarrow *contingency problem*
 - Percepts provide *new* information about current state; solution is a tree or policy; often interleave search and execution.
 - Unknown state space \Rightarrow *exploration problem* ("online")
 - When states and actions of the environment are unknown.

10

Example: vacuum world


- Single state, start in #5.
Solution??



11

Example: vacuum world

- Single state, start in #5.
Solution??
 - [Right, Suck]



12


Example: vacuum world

- **Single-state**, start in #5.
Solution? [Right, Suck]

■

- **Sensorless**, start in {1,2,3,4,5,6,7,8} e.g.,
Right goes to {2,4,6,8}
Solution?

■



Example: vacuum world


- **Sensorless**, start in {1,2,3,4,5,6,7,8} e.g.,
Right goes to {2,4,6,8}

Solution?

[Right, Suck, Left, Suck]

- **Contingency**

- Nondeterministic: Suck may dirty a clean carpet
- Partially observable: location, dirt at current location.
- Percept: [A, Clean], i.e., start in #5 or #7
Solution?




Example: vacuum world

- **Sensorless**, start in {1,2,3,4,5,6,7,8} e.g.,
Right goes to {2,4,6,8}
Solution?

[Right, Suck, Left, Suck]

- **Contingency**

- Nondeterministic: Suck may dirty a clean carpet
- Partially observable: location, dirt at current location.
- Percept: [A, Clean], i.e., start in #5 or #7
Solution? [Right, if dirt then Suck]



Single-state problem formulation


A **problem** is defined by four items:

1. **initial state** e.g., "at Arad"
 2. **actions** or **successor function** $S(x)$ = set of action-state pairs
□ e.g., $S(\text{Arad}) = \{ \langle \text{Arad} \rightarrow \text{Zerind}, \text{Zerind} \rangle, \dots \}$
 3. **goal test**, can be
□ **explicit**, e.g., $x = \text{"at Bucharest"}$
□ **implicit**, e.g., $\text{Checkmate}(x)$
 4. **path cost** (additive)
□ e.g., sum of distances, number of actions executed, etc.
□ $c(x,a,y)$ is the **step cost**, assumed to be ≥ 0
- A **solution** is a sequence of actions leading from the initial state to a goal state

Selecting a state space


- Real world is absurdly complex
→ state space must be **abstracted** for problem solving
- (Abstract) state = set of real states
- (Abstract) action = complex combination of real actions
□ e.g., "Arad → Zerind" represents a complex set of possible routes, detours, rest stops, etc.
- For guaranteed realizability, **any** real state "in Arad" must get to **some** real state "in Zerind"
- (Abstract) solution =
□ set of real paths that are solutions in the real world
- Each abstract action should be "easier" than the original problem

Vacuum world state space graph



- *states?*
- *actions?*
- *goal test?*
- *path cost?*
-

Vacuum world state space graph




- **states?** integer dirt and robot location
- **actions?** Left, Right, Suck
- **goal test?** no dirt at all locations
- **path cost?** 1 per action



19

Example: The 8-puzzle




- **states?**
- **actions?**
- **goal test?**
- **path cost?**



20

Example: The 8-puzzle



- **states?** locations of tiles
- **actions?** move blank left, right, up, down
- **goal test?** = goal state (given)
- **path cost?** 1 per move



21


Basic search algorithms

- How do we find the solutions of previous problems?
 - Search the state space (remember complexity of space depends on state representation)
 - Here: search through *explicit tree generation*
 - ROOT= initial state.
 - Nodes and leaves generated through successor function.
 - In general search generates a graph (same state through multiple paths)



22

Simple tree search example




```
function TREE-SEARCH(problem, strategy) return a solution or failure
  Initialize search tree to the initial state of the problem
  do
    if no candidates for expansion then return failure
    choose leaf node for expansion according to strategy
    if node contains goal state then return solution
    else expand the node and add resulting nodes to the search tree
  enddo
```



23

Simple tree search example




```
function TREE-SEARCH(problem, strategy) return a solution or failure
  Initialize search tree to the initial state of the problem
  do
    if no candidates for expansion then return failure
    choose leaf node for expansion according to strategy
    if node contains goal state then return solution
    else expand the node and add resulting nodes to the search tree
  enddo
```



24

Simple tree search example




```
function TREE-SEARCH(problem, strategy) return a solution or failure
  Initialize search tree to the initial state of the problem
  do
    if no candidates for expansion then return failure
    choose leaf node for expansion according to strategy ← Determines search process!!
    if node contains goal state then return solution
    else expand the node and add resulting nodes to the search tree
  enddo
```



25

State space vs. search tree

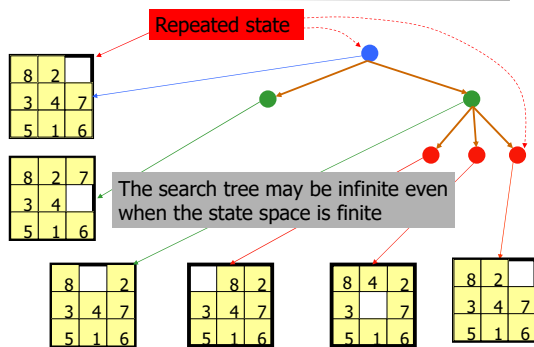


- A *state* is a (representation of) a physical configuration
- A *node* is a data structure belong to a search tree
 - A node has a parent, children, ... and includes path cost, depth, ...
 - Here *node* = *<state, parent-node, action, path-cost, depth>*
 - *FRONTIER* = contains generated nodes which are not yet expanded.
 - White nodes with black outline



26

Search Nodes \neq States



27

Tree search algorithm

```
function TREE-SEARCH(problem, frontier) return a solution or failure
  frontier ← INSERT(MAKE-NODE(INITIAL-STATE[problem]), frontier)
  loop do
    if EMPTY?(frontier) then return failure
    node ← REMOVE-FIRST(frontier)
    if GOAL-TEST[problem] applied to STATE[node] succeeds
      then return SOLUTION(node)
    frontier ← INSERT-ALL(EXPAND(node, problem), frontier)
```



28

Tree search algorithm (2)

```
function EXPAND(node, problem) return a set of nodes
  successors ← the empty set
  for each <action, result> in SUCCESSOR-FN[problem](STATE[node]) do
    s ← a new NODE
    STATE[s] ← result
    PARENT-NODE[s] ← node
    ACTION[s] ← action
    PATH-COST[s] ← PATH-COST[node] + STEP-COST(node, action, s)
    DEPTH[s] ← DEPTH[node] + 1
    add s to successors
  return successors
```



29

Search strategies

- A strategy is defined by picking the order of node expansion.
- Problem-solving performance is measured in four ways:
 - **Completeness**; Does it always find a solution if one exists?
 - **Optimality**; Does it always find the least-cost solution?
 - **Time Complexity**; Number of nodes generated/expanded?
 - **Space Complexity**; Number of nodes stored in memory during search?
- Time and space complexity are measured in terms of problem difficulty defined by:
 - *b* - maximum branching factor of the search tree
 - *d* - depth of the least-cost solution
 - *m* - maximum depth of the state space (may be ∞)



30

Blind vs. Heuristic Strategies

- **Blind** (or **uninformed**) strategies do not exploit any of the information contained in a state
- **Heuristic** (or **informed**) strategies exploits such information to assess that one node is "more promising" than another

31

Uninformed search strategies

- Categories defined by expansion algorithm:
 - ☑ **Breadth-first search**
 - ☑ **Uniform-cost search**
 - ☑ **Depth-first search**
 - ☐ Depth-limited search
 - ☐ Iterative deepening search.
 - ☐ Bidirectional search

32

Uninformed search strategies

- **Breadth-first**
 - ☐ Bidirectional
- **Depth-first**
 - ☐ Depth-limited
 - ☐ Iterative deepening


Step cost = 1

- **Uniform-Cost**
 - Step cost = $c(\text{action}) \geq \epsilon > 0$

33

Breadth-first search


- Expand shallowest unexpanded node
- **Implementation:**
 - ☐ *frontier* is a FIFO queue, i.e., new successors go at end
 - ☐



34

Breadth-first search


- Expand shallowest unexpanded node
- **Implementation:**
 - ☐ *frontier* is a FIFO queue, i.e., new successors go at end
 - ☐



35

Breadth-first search

- Expand shallowest unexpanded node
- **Implementation:**
 - ☐ *frontier* is a FIFO queue, i.e., new successors go at end
 - ☐




36

Breadth-first search

- Expand shallowest unexpanded node

Implementation:

- ☐ *frontier* is a FIFO queue, i.e., new successors go at end
- ☐



37

Properties of breadth-first search

- Complete?** Yes (if b is finite)
- Optimal?** Yes (if cost = 1 per step)
- Time?** $1 + b + b^2 + b^3 + \dots + b^d + b(b^d - 1) = O(b^{d+1})$
- Space?** $O(b^{d+1})$ (keeps every node in memory)
- Space** is the bigger problem (more than time)

38

Time and Memory Requirements

d	#Nodes	Time	Memory
2	111	.01 msec	11 Kbytes
4	11,111	1 msec	1 Mbyte
6	$\sim 10^6$	1 sec	100 Mb
8	$\sim 10^8$	100 sec	10 Gbytes
10	$\sim 10^{10}$	2.8 hours	1 Tbyte
12	$\sim 10^{12}$	11.6 days	100 Tbytes
14	$\sim 10^{14}$	3.2 years	10,000 Tb

Assumptions: $b = 10$; 1,000,000 nodes/sec; 100bytes/node

39

Time and Memory Requirements



d	#Nodes	Time	Memory
2	111	.01 msec	11 Kbytes
4	11,111	1 msec	1 Mbyte
6	$\sim 10^6$	1 sec	100 Mb
8	$\sim 10^8$	100 sec	10 Gbytes
10	$\sim 10^{10}$	2.8 hours	1 Tbyte
12	$\sim 10^{12}$	11.6 days	100 Tbytes
14	$\sim 10^{14}$	3.2 years	10,000 Tb

Assumptions: $b = 10$; 1,000,000 nodes/sec; 100bytes/node

40

Uniform-Cost Strategy

- Each step has some cost $\geq \epsilon > 0$.
- The cost of the path to each frontier node N is $g(N) = \sum$ costs of all steps.
- The goal is to generate a solution path of minimal cost.
- The queue **FRONTIER** is sorted in increasing cost.



41

Depth-first search


- Expand deepest unexpanded node
- Implementation:**
 - ☐ *frontier* = LIFO queue, i.e., put successors at front
 - ☐



42

Depth-first search

- Expand deepest unexpanded node
- Implementation:
 - ☐ frontier = LIFO queue, i.e., put successors at front
 - ☐



43

Depth-first search


- Expand deepest unexpanded node
- Implementation:
 - ☐ frontier = LIFO queue, i.e., put successors at front
 - ☐



44

Depth-first search

- Expand deepest unexpanded node
- Implementation:
 - ☐ frontier = LIFO queue, i.e., put successors at front
 - ☐



45

Depth-first search


- Expand deepest unexpanded node
- Implementation:
 - ☐ frontier = LIFO queue, i.e., put successors at front
 - ☐



46

Depth-first search


- Expand deepest unexpanded node
- Implementation:
 - ☐ frontier = LIFO queue, i.e., put successors at front
 - ☐



47

Depth-first search


- Expand deepest unexpanded node
- Implementation:
 - ☐ frontier = LIFO queue, i.e., put successors at front
 - ☐



48

Depth-first search


- Expand deepest unexpanded node
- Implementation:
 - ☐ frontier = LIFO queue, i.e., put successors at front
 - ☐



49

Depth-first search


- Expand deepest unexpanded node
- Implementation:
 - ☐ frontier = LIFO queue, i.e., put successors at front
 - ☐



50

Depth-first search


- Expand deepest unexpanded node
- Implementation:
 - ☐ frontier = LIFO queue, i.e., put successors at front
 - ☐



51

Depth-first search


- Expand deepest unexpanded node
- Implementation:
 - ☐ frontier = LIFO queue, i.e., put successors at front
 - ☐



52

Depth-first search

- Expand deepest unexpanded node
- Implementation:
 - ☐ frontier = LIFO queue, i.e., put successors at front
 - ☐



53

Properties of depth-first search

- **Complete?** No: fails in infinite-depth spaces, spaces with loops
 - ☐ Modify to avoid repeated states along path
 - complete in finite spaces
- **Optimal?** No
- **Time?** $O(b^m)$: terrible if m is much larger than d
 - ☐ but if solutions are dense, may be much faster than breadth-first
- **Space?** $O(bm)$, i.e., linear space!

54

Summary



- Search tree \neq state space
- Search strategies: breadth-first, depth-first, and variants
- Evaluation of strategies: completeness, optimality, time and space complexity
- Avoiding repeated states
- Optimal search with variable step costs



55